

NLP course 2022

Homework 1

Named Entity Recognition

Prof. Roberto Navigli

Teaching assistants:

Andrea Bacciu, Stefan Bejgu,
Abelardo Carlos Martinez Lorenzo,
Valerio Neri, Riccardo Orlando,
Alessandro Scirè, Simone Tedeschi



SAPIENZA
UNIVERSITÀ DI ROMA



Named Entity Recognition

An introduction



SAPIENZA
NLP

What is NER?

- NER is the task of **locating** and **classifying** named entities mentioned in text.
 - It is a sequence labeling task, a.k.a. every token must be tagged with some label
 - Six labels (plus **O** tag for no named entity)¹:

PER: Person

CORP: Corporation

LOC: Location

PROD: Product

GRP: Group

CW: Creative Work



¹The number of classes depends on the dataset.

Model



SAPIENZA
NLP



Model: possible approaches

You can start from the bonus exercise model!

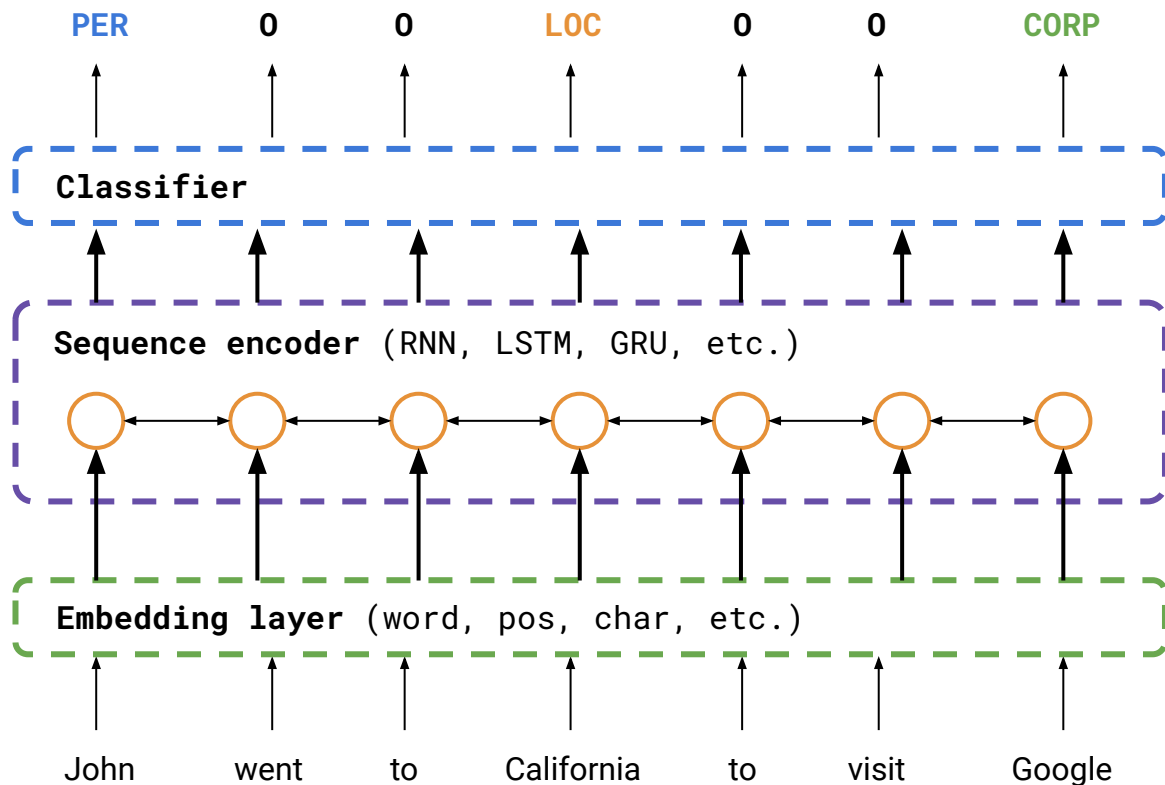
- You probably need to **encode the sequence**
 - RNNs 🙄
 - ...

Hint: Plain accuracy probably is not the best measure to use... Why?

- **Look at the data:** things are going to be clearer!

Note: for this homework, Transformer-based architectures are not allowed

Model: possible approaches



Submission

What you will receive

- We will provide you with a folder organized as follows (some files are omitted):

- nlp2022-hw1/
 - data/
 - hw1/
 - model.py
 - **stud/**
 - **model/**
 - **requirements.txt**
 - test.sh

- You are allowed to edit only the items in bold!

What you will receive

- We will evaluate your work using Docker
 - You should be fine even if you don't know anything about it
- If **test.sh** runs on your side, it will run on ours as well
 - Just keep in mind: do not change any file but those we marked in bold as editable in the previous slide
- Additionally, we wrote a **README.md** to get you everything up and running
- You can find the code repository [here](#)!

What we expect from you

- The zip folder we gave you (but populated :))
- Put your training code (if you used Colab, download the notebook `.ipynb` and place it) in `hw1/stud/`
- If you use any additional library, modify the `requirements.txt` file as needed (click [here](#) for info)
- Use the data (train and dev) in the data folder
 - use each file as defined in the **standard ML conventions** (*train for training, dev for model selection*)

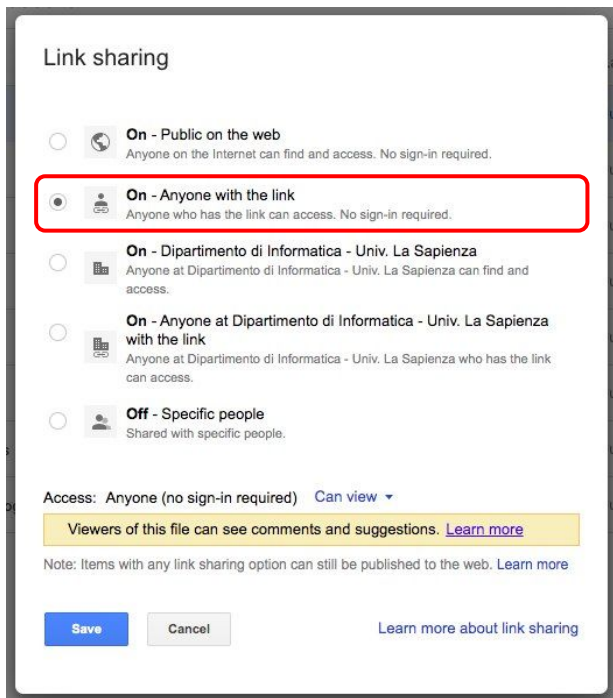
What we expect from you

- Put everything your model needs (vocabulary, weights, ...) inside the **model/** folder, and be sure to properly load them in your model
- In **hw1/stud/implementation.py** implement the **StudentModel** class
 - Load your model and use it in the **predict** method
 - You must respect the signature of the predict method!
 - You can add other methods (i.e. the constructor)
- In **hw1/stud/implementation.py** implement the **build_model** function
 - It should initialize your **StudentModel** class.

What we expect from you

- Use **test.sh** to check that everything works
- Add your **report.pdf** to the folder (yes, export it in PDF even if you are using Word!)
- Name the zip folder **lastname_studentid_hw1.zip**:
 - Ex: Luigi D'Andrea will submit a file named **dandrea_1234567_hw1.zip**
 - If you are unsure which name to put, use the one in your institutional email account

Submission Instructions



The screenshot shows the 'Link sharing' settings for a Google Drive file. The 'On - Anyone with the link' option is selected and highlighted with a red rectangle. Below it, the 'Access' is set to 'Anyone (no sign-in required)'. A yellow banner indicates that viewers can see comments and suggestions. At the bottom, there are 'Save' and 'Cancel' buttons, and a link to 'Learn more about link sharing'.

Link sharing

☐ On - Public on the web
Anyone on the Internet can find and access. No sign-in required.

☒ On - Anyone with the link
Anyone who has the link can access. No sign-in required.

☐ On - Dipartimento di Informatica - Univ. La Sapienza
Anyone at Dipartimento di Informatica - Univ. La Sapienza can find and access.

☐ On - Anyone at Dipartimento di Informatica - Univ. La Sapienza with the link
Anyone at Dipartimento di Informatica - Univ. La Sapienza who has the link can access.

☐ Off - Specific people
Shared with specific people.

Access: Anyone (no sign-in required) [Can view](#)

Viewers of this file can see comments and suggestions. [Learn more](#)

Note: Items with any link sharing option can still be published to the web. [Learn more](#)

[Save](#) [Cancel](#) [Learn more about link sharing](#)

- Upload the zip on your **institutional** Drive and make it **link-shareable** and **public** to anyone (an automatic script will download it).
- Make sure it is accessible via an incognito page of your browser!
- Do **NOT modify** the folder structure
- You have to submit the homework through the submission form on Google Classroom. You will be asked to fill a form with the requested information and the **link** to the zip you uploaded on Drive.

Dataset and Evaluation

The Dataset

- The dataset focuses on detecting **semantically ambiguous** and **complex entities** in short and low-context settings
- It contains
 - **14535** training examples
 - **765** development examples
- A label can cover a **span** of words
 - Usually it is addressed using the BIO tags
 - So there are actually **13** classes: B-PER, B-LOC, B-GRP, B-CORP, B-PROD, B-CW, I-PER, I-LOC, I-GRP, I-CORP, I-PROD, I-CW and O

What is BIO format?

BIO stands for **B**eginning, **I**nside and **O**utside of a text segment that compose a Named Entity

- **B**- prefix before a tag indicates that the tag is the beginning of a chunk
- **I**- prefix before a tag indicates that the tag is inside a chunk
- **O**- tag indicates that a token belongs to **NO** entity/chunk

PER	Minjun	B-PER
	is	O
	from	O
LOC	South	B-LOC
	Korea	I-LOC
	.	O

The dataset

```
1 # id 0
2 ease B-CW
3 on I-CW
4 down I-CW
5 the I-CW
6 road I-CW
7 - 0
8 charlie B-PER
9 smalls I-PER
10 ( 0
11 diana B-PER
12 ross I-PER
13 and 0
14 michael B-PER
15 jackson I-PER
16 in 0
17 the B-CW
18 wiz I-CW
19 ) 0
20
21 # id 1
22 it 0
23 was 0
24 republished 0
25 by 0
26 mit B-CORP
```

- The dataset is a TSV where each line is a word tab-separated with its label
 - There is a header at the beginning with the sentence ID
 - The end of a sentence is marked with an empty line

Evaluation

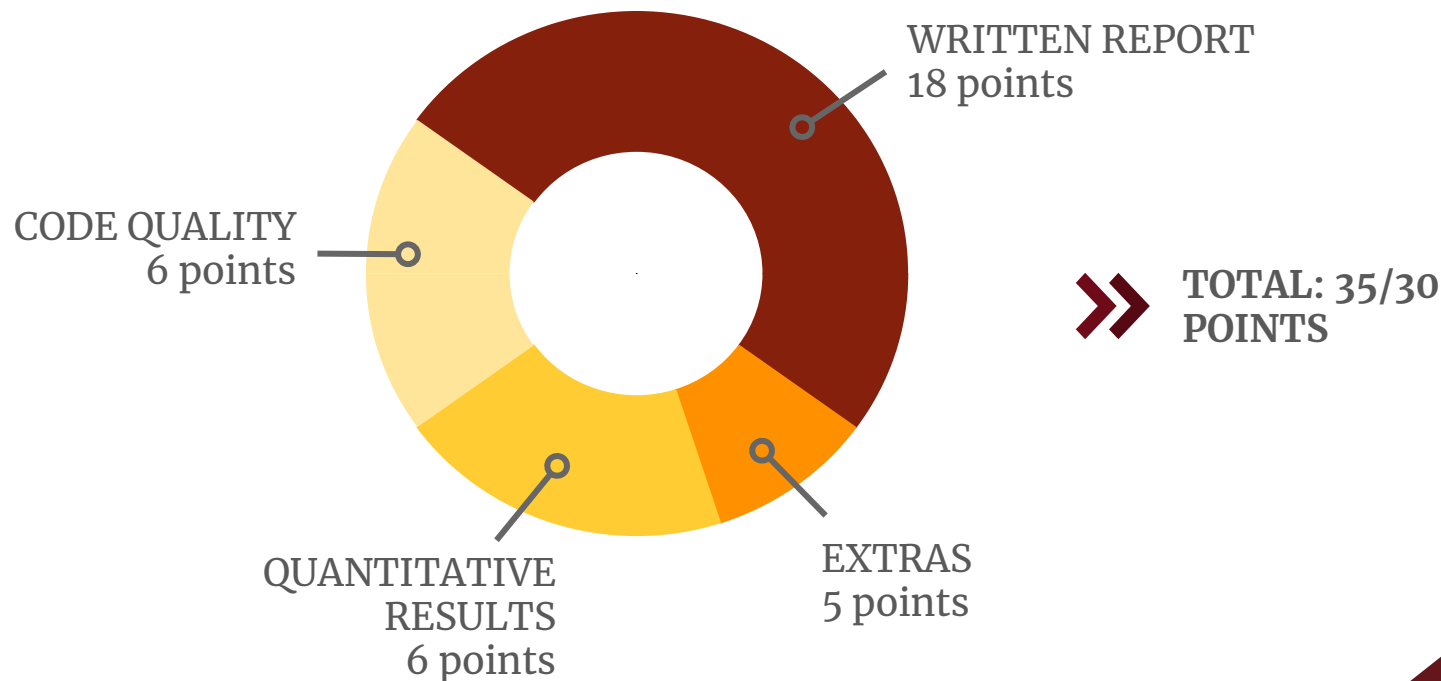
- The evaluation will be conducted on a **SECRET** test set
- We will use [segeval](#) to compute the F1 of your systems:

$$F_1 = 2 \cdot \frac{\text{recall} \times \text{precision}}{\text{recall} + \text{precision}}$$

- Use the validation split to select the best model/hyperparameters configuration

Evaluation

We will take into account the following criteria:



Report: dos and don'ts

- **ACL 2021 paper template**
 - Available [here](#) (Word and LaTeX direct download) or [here](#) (Overleaf LaTeX template)
 - You can use either the LaTeX or the Word template, your choice
 - **DO NOT MODIFY** the template (margins, spacing, font size)
 - Use the non-anonymous flag, so you can enter your name
- **Max 2 pages**
 - For the report, including title, subtitles, etc.
 - This is a **STRICT RULE!**
- **Unlimited extra pages for images, tables and references**
 - Be sure to **include** and properly **comment** a [confusion matrix](#), visualized as heat map
 - Every image and table must have a caption (don't abuse them please :))
 - Tables and images must be referenced in the report

Report: what you are expected to do



We expect a good report to be:

- **Readable and understandable**
 - We will not give penalties for English errors, but we expect the report to follow a clear flow. We don't want to read just a sequence of statements on what you did without showing the reasoning behind your choices
- **Well-structured and organized**
 - Take inspiration from the many papers available online and organize your report in well-defined sections (e.g. method, setup, experiments, results...)

Report: what you are not expected to do



We expect a good report **NOT** to include:

- Unnecessary **task** or **dataset descriptions**
 - just focus on your solution to the problem.
- **Code** copy-paste
 - Your code should be self-explanatory, so no need to show it in the report. You can add **pseudocode** to show some particular algorithm, but **no code or screenshots**, please!

Report: what you are not expected to do



We expect a good report **NOT** to include:

- **Unnecessary low-level implementation details**
 - Avoid any **low-level implementation/technical details** like “I used a dictionary to store these values”, “I had to use configuration X to solve this exception”, “I could not use Y because there was a dependency issue with Z”, etc.
 - Instead, **we are interested in high-level abstractions/strategies** you decide to use to tackle the homework, as well as the **intuitions behind your choices**.
E.g. use and description of a particular model, explanation of how and why an architecture works, etc.

Application: what you are expected to do



Your project should conform to the following rules:

- You **MUST** use PyTorch.
 - TensorFlow and other deep learning frameworks are **NOT** allowed.
 - PyTorch Lightning is **NOT** allowed (at this stage)
- **Frameworks** that use PyTorch (e.g. AllenNLP, torchtext...) are **NOT** allowed.
- Libraries (such as tqdm, sklearn, NLTK) are fine, but since the line between a framework and a library is sometimes blurred, please ask in the Google Classroom group before using any external library: **any other library MUST be agreed with the TAs.**

Application: what you are not expected to do



Your project should conform to the following rules:

- **You are not allowed** to use tools/architectures that have not been explained yet in the course, in particular:
 - word embeddings (Word2Vec, GloVe, etc.) **are allowed**,
 - contextualized word embeddings (ELMo, etc.) **are NOT allowed**,
 - Transformer-based models (BERT, BART, RoBERTa, etc.) **are NOT allowed**.
- For any doubt, please ask the TAs on Google Classroom.
- **Comment** your code, please!

Quantitative Results

We will evaluate the **performance of your model** on a SECRET test set.

You can get **from 0 to 6** points according to the following **thresholds**:

- $P < 0.15$ \Rightarrow FAIL
- $0.15 < P < 0.30$ \Rightarrow 0
- $0.30 < P < T2$ \Rightarrow 1
- $T2 < P < T3$ \Rightarrow 2
- $T3 < P < T4$ \Rightarrow 3
- $T4 < P < T5$ \Rightarrow 4
- $T5 < P < T6$ \Rightarrow 5
- $P > T6$ \Rightarrow 6

Quantitative Results

We will evaluate the **performance of your model** on a SECRET test set.

You can get **from 0 to 6** points according to the following **thresholds**:

- $P < 0.15$ \Rightarrow FAIL
- $0.15 < P < 0.30$ \Rightarrow 0
- $0.30 < P < T2$ \Rightarrow 1
- $T2 < P < T3$ \Rightarrow 2
- $T3 < P < T4$ \Rightarrow 3
- $T4 < P < T5$ \Rightarrow 4
- $T5 < P < T6$ \Rightarrow 5
- $P > T6$ \Rightarrow 6

Thresholds will be defined
based on an internal reference
model and the **normalized
distribution of YOUR scores!**

Extras

You can achieve **up to 5 points with some extras!**

An “extra” is whatever you decide to add to your model to make it better. For instance:

- use of pre-trained embeddings, PoS embeddings, N-grams, char embeddings, CRF, ...
- use of NLP best practices,
- comparative analysis of results in your report,
- informative plots in your report,
- **new ideas** (including using external resources in a clever way)

and more, according to internal baselines. Don't forget to **explain your choices** in the report! Extras that are not explained in the report will not be considered for evaluation.

Evaluation

- `test.sh` is identical to what we will be using
- **If it does not run on your side, we will not correct your homework**
- Note that, if you use **any kind of hard-coded paths**, this script won't work
- Use [paths relative](#) to the project root folder, e.g.:
 - **NO:** `/home/pincopallino/my_folder/model/weights.pt`
 - **OK:** `model/weights.pt`

Warnings

Things you should be aware of



SAPIENZA
NLP



Please be aware that

This is an **individual exercise**! Collaboration among the students is **not** allowed.

We will check for **plagiarism** both manually and automatically.

It is **not allowed** to:

- Copy from other students
- Share your code with other students
- Copy from online resources (StackOverflow, GitHub, Medium, Kaggle and so on).

However, you are allowed to use material from **external sources** as long as it is **not central** to the homework.

You are also allowed to use the **SOME** parts of the presented class notebooks. However, you **MUST** explicitly specify these parts in your code comments.

WARNING if a notebook uses torchtext or pytorch-lightning you **cannot** use it

Use of external data

- For your experiments, **use the provided data** (train and dev) in the data folder; use each file as defined in the standard ML conventions (train for training, dev for model selection).
- If you train it using external data, include it in the data folder.
- If you train it on dev set, it will be a **FAIL**.

Tips



A few tips to organize your work:

- **Start as soon as possible!**
 - Training a neural network requires time, possibly hours, depending on your hardware
- **Start small!**
 - If you don't get decent results with a very simple neural network, there is a good chance that adding other things won't make your model perform better
- **Leave some time for hyperparameter tuning!**
 - Sometimes good hyperparameter combinations can do wonders for your neural network
- Use Google [Colab](#) (free GPUs!)

Deadline

When to deliver what



Deadline

Submission date: **April 17th, 2022 (Sunday)**
23:59:59 Italian time (UTC + 1)

Submit the homework through the submission form on Google Classroom. You have to fill the form with the requested information and a link to the zip folder of the homework on Google Drive.

Late submission policy

Late Submission date: until **April 23rd, 2022 (Sunday)**
23:59:59 Italian time (UTC + 1)

1 point penalty will be applied for each day of delay, e.g.:

- A student delivers their homework on April 20th -> max possible grade $35 - 3 = 32$
- A student delivers their homework on April 23rd -> **FAIL!**

Awards

Get a **Sapienza NLP™** t-shirt



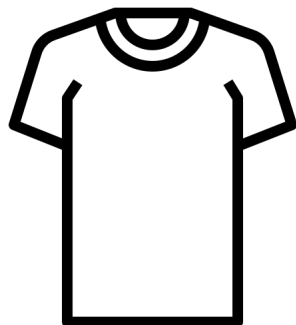
SAPIENZA
NLP



Win a Sapienza NLP t-shirt!

We will hand out amazing Sapienza NLP t-shirts to the **overall top-5** students!

The final ranking will be computed according to the scores on our **secret** test set.



That's not all

If your work is novel, interesting and original, we will gladly invite you to work together with us to extended on a fully-fledged paper for **TOP-TIER INTERNATIONAL CONFERENCE!**

Just over the last 12 months, the Sapienza NLP group published more than a dozen of papers!

Questions?

If you have a question that may interest your colleagues, **please ask it on Google Classroom.**

Otherwise, for personal or other questions, email **ALL** of us (but please, only reach for things that can't be asked on the Google Classroom).

Our emails are:

`{bacciu, bejgu, neri, orlando, scire, tedeschi}@diag.uniroma1.it`

Good Luck!!

